

# SQL for Querying and Updating

Acknowledgement:

Most of these slides have been adapted from slides made available by Pearson/Addison Wesley to accompany Elmasri and Navathe's textbook Fundamentals of Database Systems

# Outline

- The SELECT Structure of SQL Queries
  - Aliasing, and Renaming
  - Joined Tables
  - Aggregate Functions and Grouping
  - Tables as Sets in SQL
  - Nested Queries
  - Dealing with NULLs
  - The CASE Clause, and Result Limitation
- DELETE, and UPDATE Statements

# Retrieval Queries in SQL

- **SELECT statement**
  - One basic statement for retrieving information from a database
- **SQL allows a table to have two or more tuples that are identical in all their attribute values**
  - Unlike relational model (relational model is strictly set-theory based)
  - Multiset or bag behavior
  - Tuple-id may be used as a key

# The SELECT Structure of SQL Queries

## ■ Basic form of the SELECT statement:

Relational algebra correspondence:

|                                             |   |                                                      |
|---------------------------------------------|---|------------------------------------------------------|
| <b>SELECT</b> <attribute and function list> | → | Project ( $\pi$ )                                    |
| <b>FROM</b> <table list>                    | → | Cartesian Product ( $\times$ ) or Join ( $\bowtie$ ) |
| [ <b>WHERE</b> <condition> ]                | → | Select ( $\sigma$ )                                  |
| [ <b>GROUP BY</b> <grouping attribute(s)> ] | → | Aggregate/Grouping ( $\gamma$ )                      |
| [ <b>HAVING</b> <group condition> ]         |   |                                                      |
| [ <b>ORDER BY</b> <attribute list> ];       |   |                                                      |

where

- <attribute list> is a list of attribute names whose values are to be retrieved by the query.
- <table list> is a list of the relation names required to process the query.
- <condition> is a conditional (Boolean) expression that identifies the tuples to be retrieved by the query.

# The SELECT Structure of SQL Queries (2)

- Projection attributes
  - Attributes whose values are to be retrieved
  - Specify an asterisk (\*) to retrieve all the attribute values
  - The projection can include functions and formulas
- Selection condition
  - Boolean condition that must be true for any retrieved tuple
  - Selection conditions include join conditions when multiple relations are involved

# Logical and Arithmetic Operators, and Substring Pattern Matching

- Logical comparison operators: =, <, <=, >, >=, <>, etc.
- **BETWEEN** comparison operator
  - Example: **WHERE** Salary **BETWEEN** 30000 **AND** 40000
- **LIKE** comparison operator
  - Used for string **pattern matching**
  - % replaces an arbitrary number of zero or more characters
  - underscore (\_) replaces a single character
    - Examples: **WHERE** Address **LIKE** '%Houston,TX%';  
**WHERE** Ssn **LIKE** '\_\_1\_\_8901';
- Standard arithmetic operators such as (+, −, \*, /) and functions may be included as a part of **SELECT**
  - Example: **SELECT** salary \* 0.9 **FROM** EMPLOYEE;  
**WHERE** age(Bdate) >= 21;

# Ambiguous Attribute Names

- Same name can be used for two (or more) attributes in different relations
  - As long as the attributes are in different relations
  - Must **qualify** the attribute name with the relation name to prevent ambiguity
  - The \* can be prefixed by the relation name to retrieve all attributes of a table

**Example:** Retrieve all attributes of the manager of the “Research” department

```
SELECT EMPLOYEE.*
```

```
FROM EMPLOYEE, DEPARTMENT
```

```
WHERE DEPARTMENT.Name='Research' AND
```

```
DEPARTMENT.Dnumber=EMPLOYEE.Dnumber;
```

# Aliasing, and Renaming

- Aliases or tuple variables
  - Are needed to refer to the same table more than once in a query
  - Can also be used just to shorten the query
- Renaming of attributes
  - Are needed to distinguish attributes that are homonyms
  - Can also be used for convenience

**Example:** For each employee, retrieve the employee's first and last name and the first and last name of his or her immediate supervisor

```
SELECT E.Fname, E.Lname,  
       S.Fname AS Sup_Fname, S.Sname AS Sup_Sname  
FROM EMPLOYEE AS E, EMPLOYEE AS S  
WHERE E.Super_ssn=S.Ssn;
```



# Aliasing, and Renaming (2)

- The attribute names can also be renamed in the tuple variable

**Example:**

```
SELECT ...
```

```
FROM EMPLOYEE AS E(Fn, Mi, Ln, Ssn, Bd, Addr, Sex, Sal,  
Sssn, Dno)...
```

- The “AS” may be dropped in most SQL implementations

# Joined Tables

- The join operation can be represented by:
  - A cartesian product followed by a selection, or
  - A joined table
    - Specify different types of join
      - [INNER] JOIN, NATURAL JOIN, LEFT|RIGHT|FULL [OUTER] JOIN
- Example: Retrieve all attributes of the manager of the “Research” department

```
SELECT e.*
FROM employee AS e JOIN department AS d
ON e.dno=d.dnumber
WHERE d.name='Research';
```

# Joined Tables (2)

- A joined table can nest join specifications
  - Optionally enclosed in parentheses
  - Example: Retrieve the first and last names of all employees and, for those who work on projects, the name of the project and the number of hours per week

```
SELECT e.fname, e.lname, p.pname, w.hours
FROM employee AS e
LEFT OUTER JOIN works_on AS w ON e.ssn=w.essn
LEFT OUTER JOIN project AS p ON w.pno=p.pnumber;
```

# Aggregate Functions in SQL

- Used to summarize information from multiple tuples into a single-tuple summary
- Built-in aggregate functions
  - COUNT, SUM, MAX, MIN, AVG, STDDEV\_POP, ...
  - NULLs are discarded
- Grouping
  - Create subgroups of tuples before summarizing
- To select entire groups, HAVING clause is used
  - HAVING corresponds to a posterior select operation
- Aggregate functions can be used in the SELECT clause or in a HAVING clause

# Aggregate Functions in SQL (2)

- Examples

- Retrieve the number of employees, the number of supervisees and the number supervisors

```
SELECT COUNT(*) AS num_employees,  
       COUNT(super_ssn) AS num_supervisees,  
       COUNT(DISTINCT(super_ssn)) AS num_supervisors  
FROM employee;
```

- Retrieve the average salary for employees per sex

```
SELECT sex, AVG(salary) AS avg_sal  
FROM employee  
GROUP BY sex;
```

- Retrieve the names of the departments with at least 3 employees

```
SELECT d.dname  
FROM department d JOIN employee e ON d.dnumber=e.dno  
GROUP BY d.dname  
HAVING COUNT(*) >= 3;
```

# Aggregate Functions in SQL (3)

- Pitfalls when combining the WHERE and the HAVING Clause
  - Example: Retrieve the total number of employees whose salaries exceed \$30,000 in each department, but only for departments where at least 3 employees work
  - **Incorrect** query:

```
SELECT dno, COUNT(*)  
FROM employee  
WHERE salary > 30000  
GROUP BY dno  
HAVING COUNT(*) >= 3;
```
  - **Correct** query:

```
SELECT dno, COUNT(*)  
FROM employee  
WHERE salary > 30000  
AND dno IN (SELECT dno FROM employee  
            GROUP BY dno HAVING COUNT(*) >= 3)  
GROUP BY dno;
```

# Ordering of Query Results

- Use **ORDER BY** clause
  - Keyword **DESC** to see result in a descending order of values
  - Keyword **ASC** to specify ascending order explicitly
  - Typically placed at the end of the query

SELECT ...

ORDER BY D.Dname DESC, E.Lname ASC,  
E.Fname ASC;

# Tables as Sets in SQL

- SQL does not automatically eliminate duplicate tuples in query results
- Use the keyword **DISTINCT** in the SELECT clause
  - Only distinct tuples should remain in the result

## Example:

```
SELECT DISTINCT E.SSN  
FROM EMPLOYEE E, DEPENDENT D  
WHERE E.SSN=D.ESSN;
```



# Tables as Sets in SQL (2)

- Set operations
  - UNION, EXCEPT (difference), INTERSECT
    - Corresponding multiset operations: UNION ALL, EXCEPT ALL, INTERSECT ALL
  - Type compatibility is needed for these operations to be valid

**Query 4.** Make a list of all project numbers for projects that involve an employee whose last name is 'Smith', either as a worker or as a manager of the department that controls the project.

```
Q4A: ( SELECT   DISTINCT Pnumber
      FROM      PROJECT, DEPARTMENT, EMPLOYEE
      WHERE     Dnum=Dnumber AND Mgr_ssn=Ssn
              AND Lname='Smith' )

      UNION

      ( SELECT   DISTINCT Pnumber
      FROM      PROJECT, WORKS_ON, EMPLOYEE
      WHERE     Pnumber=Pno AND Essn=Ssn
              AND Lname='Smith' );
```

# Nested Queries

- Nested queries are complete “select-from-where-group by-having-order by” blocks within WHERE clause of another query (called the outer query)
- Comparison operators
  - IN ( $v \text{ IN } S$ ): evaluates to TRUE if  $v$  is one of the elements in the set  $S$ 
    - $S$  is usually the result of a nested query
  - ANY or SOME ( $v \theta \text{ ANY } S$  (or  $\text{SOME}$ )): returns TRUE if the value  $v$  is  $\theta$  to some value in the set  $S$ 
    - $\theta$  can be any of  $>$ ,  $>=$ ,  $<$ ,  $<=$ , and  $<>$
    - When  $\theta$  is  $=$ , ANY/SOME behaves like the IN operator
  - ALL ( $v \theta \text{ ALL } S$ ): returns TRUE if the value is  $\theta$  of all values from  $S$
  - EXISTS: evaluates to true if the result of the nested query is not empty
    - “For all” queries must use NOT EXISTS

# Uncorrelated Nested Queries

- The nested query is evaluated only once
- Examples:
  - Retrieve the employees whose salary is greater than every salary from an employee who works for department 5

```
SELECT * FROM employee
WHERE salary > ALL (
  SELECT salary FROM employee
  WHERE dno=5);
```

- Retrieve all projects that do not have any woman working on them

```
SELECT * FROM project
WHERE pnumber NOT IN (
  SELECT w.pno
  FROM employee e JOIN works_on w ON e.ssn=w.essn
  WHERE e.sex='F');
```

# Correlated Nested Queries

- The nested query is evaluated once for each tuple in the outer query
- Examples:
  - Retrieve the employees who have no female dependent

```
SELECT * FROM employee e
WHERE NOT EXISTS (
  SELECT * FROM dependent d
  WHERE e.ssn=d.essn AND d.sex='F');
```

# Representing the Relational Division using NOT EXISTS

- Example: Retrieve the names of employees who work on all the projects controlled by department 5

```
SELECT e.fname, e.minit, e.lname FROM employee e
WHERE NOT EXISTS (
  (SELECT p.pnumber FROM project p WHERE p.dnum=5)
  EXCEPT
  (SELECT w.pno FROM works_on w WHERE e.ssn= w.essn)
);
```

```
SELECT e.fname, e.minit, e.lname FROM employee e
WHERE NOT EXISTS (
  SELECT * FROM project p
  WHERE p.dnum=5 AND NOT EXISTS
    (SELECT * FROM works_on w WHERE w.essn=e.ssn AND w.pno=p.pnumber)
);
```

# Comparisons Involving NULL and Three-Valued Logic

- Meanings of NULL
  - Unknown value
  - Unavailable or withheld value
  - Not applicable attribute
- Each individual NULL value considered to be different from every other NULL value
- SQL uses a three-valued logic:
  - TRUE, FALSE, and UNKNOWN (like Maybe)

# Comparisons Involving NULL and Three-Valued Logic (2)

- Operations involving NULL return NULL
  - Examples:
    - `NULL + 1` results in NULL
    - `(balance + overdraft_limit)` can be NULL
- Comparisons with NULL return UNKNOWN
  - Examples:
    - `SELECT * FROM employee WHERE super_ssn=NULL;`
      - Returns an empty result
    - `SELECT * FROM employee WHERE sex<>'M';`
      - Does not include in the result employees whose sex is either 'M' or NULL
- Check whether an attribute value is NULL (IS NULL or IS NOT NULL)
  - `SELECT * FROM employee WHERE super_ssn IS NULL;`
    - Returns the employees who do not have a direct supervisor

# Comparisons Involving NULL and Three-Valued Logic (3)

**Table 7.1** Logical Connectives in Three-Valued Logic

|     |            |         |         |         |
|-----|------------|---------|---------|---------|
| (a) | <b>AND</b> | TRUE    | FALSE   | UNKNOWN |
|     | TRUE       | TRUE    | FALSE   | UNKNOWN |
|     | FALSE      | FALSE   | FALSE   | FALSE   |
|     | UNKNOWN    | UNKNOWN | FALSE   | UNKNOWN |
| (b) | <b>OR</b>  | TRUE    | FALSE   | UNKNOWN |
|     | TRUE       | TRUE    | TRUE    | TRUE    |
|     | FALSE      | TRUE    | FALSE   | UNKNOWN |
|     | UNKNOWN    | TRUE    | UNKNOWN | UNKNOWN |
| (c) | <b>NOT</b> |         |         |         |
|     | TRUE       | FALSE   |         |         |
|     | FALSE      | TRUE    |         |         |
|     | UNKNOWN    | UNKNOWN |         |         |



# Useful Clauses involving NULLs

- NULLIF

- Returns NULL if the expressions are equal

- Example:

```
SELECT AVG(NULLIF(salary, 0.00)) AS avg_sal  
FROM employee;
```

- COALESCE

- Returns the first non-NULL value in a list

- Example:

```
SELECT COALESCE(update_time, create_time,  
                'Unknown') AS last_updated  
FROM file_table;
```

# The CASE Statement

- The CASE statement allows to define conditional instructions in SQL queries
- Syntax:

-- simple CASE

```
CASE ("column_name")  
  WHEN "value1" THEN "result1"  
  WHEN "value2" THEN "result2"  
  ...  
  [ELSE "resultN"]  
END
```

-- searched CASE

```
SELECT CASE  
  WHEN "condition1" THEN "result1"  
  WHEN "condition2" THEN "result2"  
  ...  
  [ELSE "resultN"]  
END
```

# The CASE Statement (2)

- Example of CASE in queries:

```
SELECT ssn, fname, lname,  
       CASE (sex)  
         WHEN 'M' THEN 'Male'  
         WHEN 'F' THEN 'Female'  
         ELSE 'Other'  
       END,  
       CASE  
         WHEN age(dt_nasc) >= 65 THEN 'Elder'  
         WHEN age(dt_nasc) < 18 THEN 'Youth'  
         ELSE 'Adult'  
       END AS age_class  
FROM employee;
```

# The CASE Statement (3)

- Example of CASE in updates:

```
UPDATE employee
SET salary = (
CASE
    WHEN dno=5 THEN salary*1.15
    WHEN dno=4 THEN salary*1.1
    ELSE salary*1.05
END);
```

# The CASE Statement (4)

- Example of CASE in aggregates:

```
SELECT dno,  
       COUNT(  
         CASE WHEN salary>30000 THEN 1  
              ELSE NULL  
         END) AS num_greater_30000,  
       COUNT(  
         CASE WHEN salary<=30000 THEN 1  
              ELSE NULL  
         END) AS num_upto_30000  
FROM employee  
GROUP BY dno;
```

# Limiting the Size of the Result

- Useful for top-queries and pagination
  - SQL:2008: FETCH FIRST|NEXT k [PERCENT] ROWS + OFFSET m [PERCENT] ROWS
  - Proprietary extensions: TOP k, [OFFSET m] LIMIT k, etc.
- Example:
  - first page  
SELECT ssn, fname, lname, salary  
FROM employee  
ORDER BY salary DESC  
FETCH FIRST 10 ROWS ONLY; -- WITH TIES
  - second page  
SELECT ssn, fname, lname, salary  
FROM employee  
ORDER BY salary DESC  
OFFSET 10 ROWS  
FETCH NEXT 10 ROWS ONLY; -- WITH TIES

# DELETE, and UPDATE Statements in SQL

- In addition to INSERT, these are essential commands used to modify the database
  - These operations are executed in transactions, either explicitly stated by the user or implicitly done by the DBMS
    - BEGIN | START TRANSACTION;  
commands;  
COMMIT | ROLLBACK;
- UPDATE may update a number of tuples (rows) in a relation (table) that satisfy the condition
- DELETE may also update a number of tuples (rows) in a relation (table) that satisfy the condition

# DELETE

- Removes tuples from a relation
  - Includes a WHERE-clause to select the tuples to be deleted
    - The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause
    - A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table
  - Referential integrity should be enforced
  - Tuples are deleted from only *one table* at a time
    - Unless CASCADE is specified on a referential integrity constraint



# DELETE (2)

- Examples:

- Delete the employees whose last name is Brown

`DELETE FROM Employee WHERE Lname='Brown';`

- Delete all employees who work for the Administration department

`DELETE FROM Employee WHERE Dno IN (  
SELECT Dnumber FROM Department  
WHERE Dname='Administration');`

- Delete all dependents

`DELETE FROM Dependent;`

# UPDATE

- Used to modify attribute values of one or more selected tuples
  - A SET-clause specifies the attributes to be modified and their new values
  - A WHERE-clause selects the tuples to be modified
- Each command modifies tuples *in the same relation*
- All constraints should be enforced

# UPDATE (2)

- Examples:
  - Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively

```
UPDATE Project
```

```
SET Plocation = 'Bellaire', DNUM = 5
```

```
WHERE Pnumber = 10;
```

- Give all employees in the department 5 a 10% raise in salary

```
UPDATE Employee
```

```
SET SALARY = SALARY*1.1
```

```
WHERE DNO = 5;
```

# The RETURNING clause

- Update statements (INSERT, DELETE, UPDATE) can have a RETURNING clause, which returns the indicated attributes of the affected tuples
- Examples
  - Insert a project returning the identifier automatically generated using a sequence

```
INSERT INTO project(pname, plocation, dnum)
VALUES('Project XYZ', 'Stafford', 5)
RETURNING pnumber;
```
  - Give all employees in the department 5 a 10% raise in salary and return their names and new salaries

```
UPDATE Employee
SET SALARY = SALARY*1.1
WHERE DNO = 5 RETURNING fname, minit, lname, salary;
```